

Philosophy of Artificial Intelligence

UCLA PHILOS 171

Summer Session C, 2026

Lectures: ¹	Paul Talma	Office Hours:
TR: 10:45am–12:35pm	ptalma@humnet.ucla.edu	TR: 12:35–1:35pm

Course Description

This course explores foundational questions in the philosophy of artificial intelligence. We will tackle the topic in three stages.

(1) We will begin by considering the origins and ambitions of AI as a field of research: what is intelligence and what would it mean for a machine to possess it? We will examine the Computational Theory of Mind, which holds that thinking is a form of computing, as well as its implementation in symbolic and connectionist architectures. Our aim here will be to acquire a solid theoretical backdrop against which to think about AI.

(2) Next, we'll examine the cognitive capacities of current AI models: can they perceive, reason, and plan? Do their words have meaning, and can they communicate? Do they have goals of their own, and can they infer the goals of others?

(3) Armed with an understanding of the *capabilities* of current AI systems, we'll consider the social and political questions they raise. What risks do generative AI models pose? What promises do they hold? How will they impact education, economic relations, the production of art and media, and the distribution of political power?

You will come away from the course equipped to think through the complex and pressing questions raised by modern AI systems. You will also refine important philosophical skills: how to analyze an argument, how to defend a position, and how to communicate with clarity and precision.

The course is self-contained: it does not presuppose any technical knowledge, and all relevant information about AI will be introduced as we go along.

¹All times Pacific.

Course Materials

Required and recommended materials will be provided on BruinLearn.

Assessment

Writing exercises: In weeks 2–6, you will have the opportunity to write a response to one of that week’s readings. You must take four of these opportunities. Each reading response consists of two paragraphs: one stating and explaining a claim from the paper, the other containing your response to the claim. Your response can take any form (constructive, critical, amplificatory, ...), so long as it engages with the claim. Responses will be graded for good-faith completion. Reading responses for each week will be due before that week’s second class.

Final paper: 6–8 pages. You will be able to choose from among several prompts, or to answer a prompt of your own design upon consultation with the instructor. You will carefully reconstruct and evaluate a philosophical position discussed in the course. Expectations for the paper will be circulated during the session.

Attendance: Your attendance in lecture is required. You will only count as present if your video is turned on for the duration of the lecture (reasonable exceptions will be made). You will be granted one excused absence (in addition to any attested medical or family emergencies).

Participation: can take many forms—asking a question in lecture, sending a follow up email, coming to office hours, etc. Doing any of these once will earn the full 5%.

Writing exercises	40%
Final paper	35%
Attendance	20%
Participation	5%

Course Modality

The course will be taught online. Attendance in lecture is required and will count toward your grade. Exceptional circumstances apart, you will need to turn on your

video to count as present. Doing so helps me gauge your reception of the material and helps you engage with the lecture. As such, it is not only to your benefit but also to that of your classmates.

Schedule

Week 1: Origins and Aims of AI

Lecture 1:

- Turing (1950), “Computing Machinery and Intelligence” (sections 1–3, 7).
- Block (1981), “Psychologism and Behaviorism.”

Lecture 2:

- Russell and Norvig (2021), *Artificial Intelligence, A Modern Approach*, Chapters 1 and 2.

Optional Readings:

- Newell and Simon (1976), “Computer Science as Empirical Enquiry: Symbols and Search.”

Week 2: The Classical Computational Theory of Mind

Lecture 1:

- Putnam (1967), “The Nature of Mental States.”
- Fodor (1987), *Psychosemantics* (Preface (ix–xiii); 10–20; Appendix (135–154)).

Lecture 2:

- Searle (1980), “Minds, Brains, and Programs.”

[Writing exercise due Friday at 11:30pm.]

Optional Readings:

- Rescorla (2024), “The Computational Theory of Mind” (Stanford Encyclopedia of Philosophy).

- Marr (1982), *Vision* (excerpts).

Week 3: The Connectionist Challenge

Lecture 1:

- Fodor and Pylyshyn (1988), “Connectionism and Cognitive Architecture” (sections 1–3, Conclusion).

Lecture 2:

- Lake and Baroni (2018), “Generalization without systematicity: On the compositional skills of sequence-to-sequence recurrent networks.”
- Goodale and Mascarenhas (2026), “Fodor and Pylyshyn’s Systematicity Challenge Still Stands.”

[Writing exercise due Friday at 11:30pm.]

Optional Readings:

- Millièrre and Buckner (2024), “A Philosophical Introduction to Language Models, Part 1.”
- LeCun et al. (2015), “Deep Learning.”
- Smolensky (1988), “On the Proper Treatment of Connectionism.”

Week 4: Functions, Goals, and Agency

Lecture 1: no reading. More on connectionism. Introduction to LLMs.

Lecture 2:

- Millikan (1993), “What is Behavior?”
- Butlin (2024), “Reinforcement Learning and Artificial Agency.”

[Writing exercise due Friday at 11:30pm.]

Optional Readings:

- Dennett (1971), “Intentional Systems.”
- Kelley (2025), “A Control Theory of Action.”

- Dung (2025), “Understanding Artificial Agency.”

Week 5: Alignment

[If writing your final paper on your own topic, we must have met by now.]

Lecture 1:

- Gabriel (2020), “Artificial Intelligence, Values, and Alignment.”

Lecture 2:

- Millièrè (2025), “Normative Conflicts and Shallow AI Alignment” (main text only).

[Writing exercise due Friday at 11:30pm.]

Optional Readings:

- Zhi-Xuan (2025), “Beyond Preferences in AI Alignment.”
- Narayanan & Kapoor (2025), “AI as Normal Technology.”
- Alexander (2025), “AI as Profoundly Abnormal Technology.”
- Lazar (2025), “Anticipatory AI Ethics.”

Week 6: Power and Property

Lecture 1:

- Winner (1980), “Do Artifacts Have Politics?”
- Lazar (2022), “Power and AI: Nature and Justification.”

Lecture 2:

- *No reading. Catch up/review.*

[Writing exercise due Friday at 11:30pm.]

Optional Readings:

- Lazar (2025), “Governing the Algorithmic City.”
- Creel and Hellman (2022), “The Algorithmic Leviathan: Arbitrariness, Fairness, and Opportunity in Algorithmic Decision-Making Systems.”
- Goetze (2024), “AI Art is Theft.”

SUNDAY 13th OF SEPTEMBER, 11:30pm: Final Paper Due

Usage of Generative AI

Writing is the most difficult and rewarding part of philosophy. It is only by wrangling your ideas down into text that you can fully understand and examine them. The arduous trial-and-error process involved in accurately, precisely, and concisely articulating a philosophical claim is an integral part of doing philosophy. It is not a mere “implementation detail,” the kind of work that gets done once the real thinking is over. It is part and parcel of the thinking itself. It is also an excellent way to learn how to write well.

For this reason, the use of generative AI to write any of the assignments is forbidden and constitutes grounds for failing the course. This includes drafting an outline and having an AI generate a paper on the basis of this outline. It also includes using an AI to write or rewrite a sentence or a paragraph. It also includes using an AI to summarize a course reading and quoting or paraphrasing the AI summary. If you are suspected of using generative AI in any of these ways, you will be required to redo the assignment under direct supervision or to discuss your work in an in-person interview. Note that current AI detection tools are extremely reliable.

You are allowed to use generative AI for the purposes of brainstorming, exploration, and checking your understanding. Doing so can be a valuable use of your time, if done properly (for example, the model needs to be prompt-engineered to give adequate pushback on your ideas). Any use of generative AI for brainstorming and the like should be cited as such.

Academic Integrity

All students are expected to adhere to the university’s academic honesty policy. Plagiarism, unauthorized collaboration, or misrepresentation of sources will result in disciplinary action. See previous section for acceptable uses of generative AI.

Accessibility

Students needing an academic accommodation based on a disability should either contact me or the Office for the Center for Accessible Education located at (310)

825-1501 or A255 Murphy Hall. When possible, students should contact the CAE within the first two weeks of the term as reasonable notice is needed to coordinate accommodations. For more information visit www.cae.ucla.edu.

Acknowledgements

This syllabus was partly inspired by those of Jocelyn Maclure, Soren Schlassa, and Dan Williams.